

# **High-Performance Machine Learning Using Scikit-Learn**

Qiyang Hu

UCLA Office of Advanced Research Computing

Aug 4<sup>th</sup>, 2023

# Outline

- Learning Scikit-learn basics
  - High-level overview of scikit-learn libraries
  - Practical usage of scikit-learn
    - Deep learning and scikit-learn
    - Scikit-learn extension libraries
- High-performance machine learning using scikit-learn-*ish tools*
  - Overview of performance issues in machine learning
  - General performance tips and tricks in scikit-learn
  - Making the computation faster
  - Processing large datasets

# Outline

- Learning Scikit-learn basics
  - High-level overview of scikit-learn libraries
  - Practical usage of scikit-learn
    - Deep learning and scikit-learn
    - Scikit-learn extension libraries
- High-performance machine learning using scikit-learn-*ish tools*
  - Overview of performance issues in machine learning
  - General performance tips and tricks in scikit-learn
  - Making the computation faster
  - Processing large datasets

# Measures of Machine Learning Performance



## Training Throughput

Number of training instances goes through the training processed in unit time




## Prediction Latency

Time to make a single prediction taken by a deployed model



## Prediction Throughput

Number of predictions made in unit time



# Factors Affecting Machine Learning Performance

## DATA

- Number of Instances
  - For training
  - For prediction
  - Upfront availability
- Features
  - Number of features
  - Importance of features
  - Feature representation



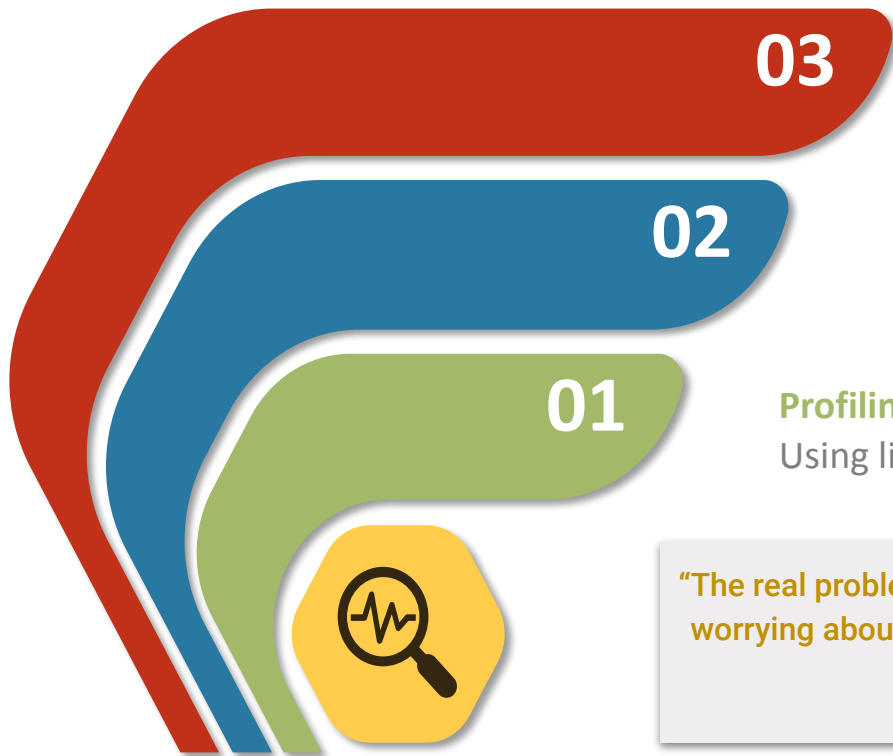
## MODEL

- Algorithm complexity
  - Hyperparameters
  - Optimized libraries
  - Multi-core Parallelism
  - Out-of-core learning
- Data complexity
  - Feature selection
  - Feature extraction
  - Feature transformation
- For prediction:
  - Validation overhead
  - Model compression
  - Model reshaping

# Outline

- Learning Scikit-learn basics
  - High-level overview of scikit-learn libraries
  - Practical usage of scikit-learn
    - Deep learning and scikit-learn
    - Scikit-learn extension libraries
- High-performance machine learning using scikit-learn *-ish tools*
  - Overview of performance issues in machine learning
  - General performance tips and tricks in scikit-learn
  - Making the computation faster
  - Processing large datasets

# Basic performance tips for scikit-learn projects



## Writing Cython wrappers

- e.g. liblinear, libsvm in LogReg and SVM models.
- Using OpenMP functions through Cython

## Using Numpy and Scipy

- Replacing nested for loops
- Performance tips for numpy and scipy

## Profiling your Python code

Using line\_profiler, memory\_profiler

“The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times...”

— Donald Knuth in “TAOCP”

## Limiting working memory

```
from sklearn import config_context as cnftxt
with cnftxt(working_memory=128):
    pass
```

## Model compression

```
clf = SGDRegressor(penalty='elasticnet', l1_ratio=0.25)
clf.fit(X_train, y_train).sparsify()
clf.predict(X_test)
```

## Configuring for reduced validation overhead

```
export SKLEARN_ASSUME_FINITE="TRUE"
```

```
from sklearn import config_context as cnftxt
with cnftxt(assume_finite=True):
    pass
```

## Scikit-Learn Tweaks & Tricks

## Model reshaping

- Selecting only a portion of the available features to fit a model
- Needs to be performed manually

## Prediction in bulk mode


Doing predictions many instances at the same time



# Outline

- Learning Scikit-learn basics
  - High-level overview of scikit-learn libraries
  - Practical usage of scikit-learn
    - Deep learning and scikit-learn
    - Scikit-learn extension libraries
- High-performance machine learning using scikit-learn-*ish tools*
  - Overview of performance issues in machine learning
  - General performance tips and tricks in scikit-learn
  - Making the computation faster
  - Processing large datasets

# Multicore Parallelism and Optimizations for Scikit-Learn



# Joblib: lightweight pipelining tools in Python



- Speeding up long-running jobs:
  - Specific optimizations for Numpy arrays
  - Caching and lazy evaluation to avoid computing repeatedly
  - No need to change the code or control flow
  - Various parallel backends: `locky`, `multiprocessing`, `threading`, `dask`, `ray`, ...
- Tight and powerful integration with Scikit-learn
  - Already used in some sklearn classes (ElasticNet, SGDClassifier, ...)
  - Simply tweak `n_jobs`: `n_jobs=-1` to use all available cores
  - Typical tasks: cross-validation, grid search, multi-label prediction, ensemble learning
- Distributed Scikit-Learn using Dask as backend

```
from dask.distributed import Client
import joblib

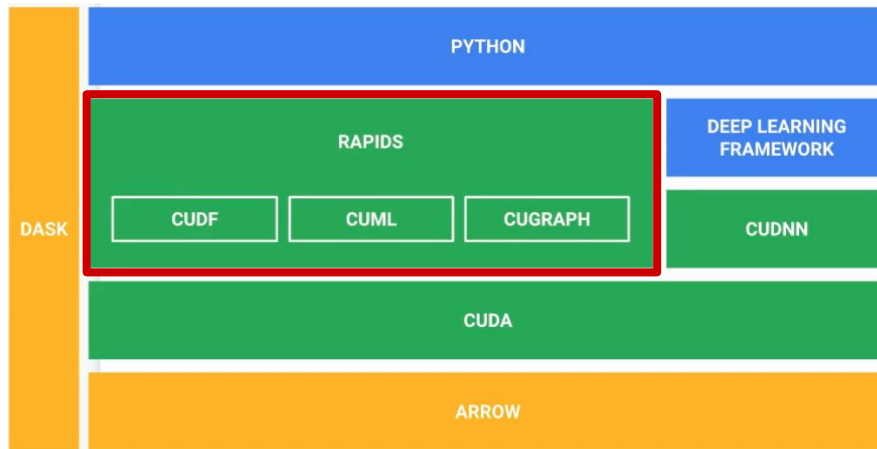
client = Client(processes=False)           # create local cluster
# client = Client("scheduler-address:8786") # or connect to remote cluster

with joblib.parallel_backend('dask'):
    # Your scikit-learn code
```

**bit.ly/hpsk1\_01**

# High-performance ML using accelerators

- Scikit-learn alone doesn't have support to GPU or TPU.
- Option 1: [RAPIDS](#) software libraries from nvidia
  - Exactly same APIs
  - Kernels rewritten by CUDA
    - *Numpy* -> *CuPy*
    - *pandas* -> *cuDF*
    - *scikit-learn* -> *cuML*
    - *networkx* -> *cuGraph*
  - Dask + RAPIDS + blazingSQL, RAPIDS + Spark, xgboost, ...



# High-performance ML using accelerators (cont'd)

- Option 2: using [scikit-learn-intelex](#)

- Higher-level APIs to daal4py
- Full conformance with all sklearn apis/algos.
- Acceleration without code change:

- Intel CPU optimization patching

```
from sklearnx import patch_sklearn
patch_sklearn()
```


- Intel CPU/GPU optimizations patching

```
from sklearnx import patch_sklearn, config_context
patch_sklearn()
with config_context(target_offload="gpu:0"):
```


- Installation:

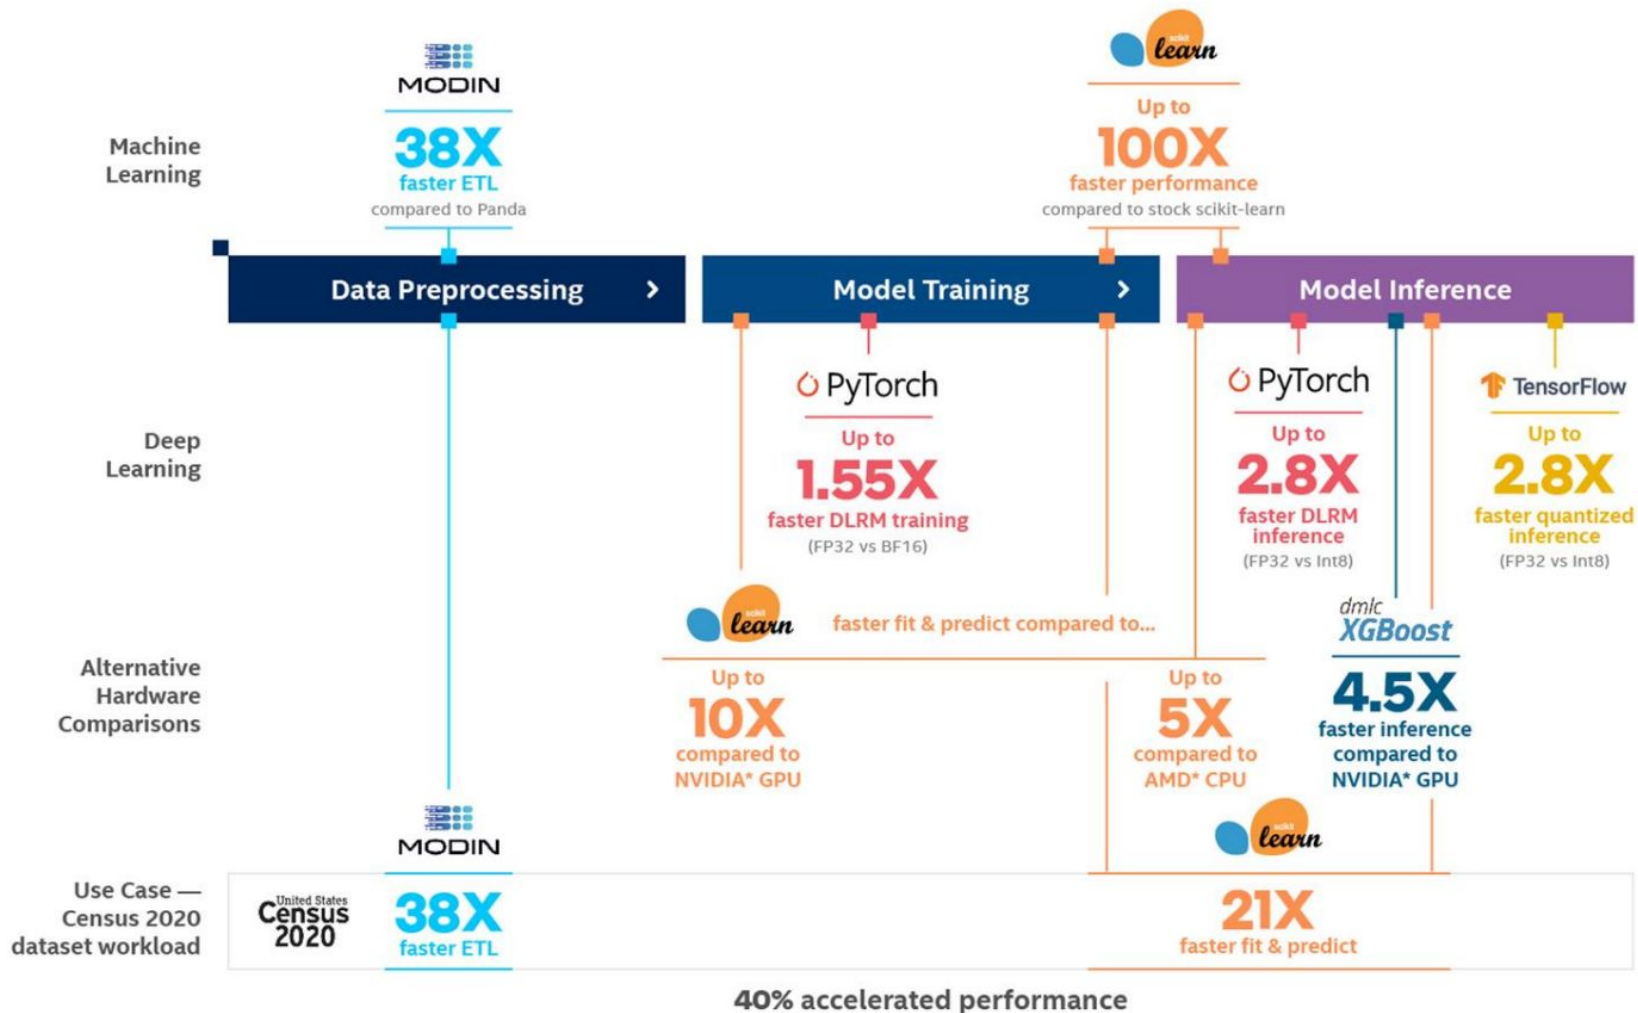
```
pip install scikit-learn-intelex
pip install dpcpp-cpp-rt
```

Speedups of Intel® Extension for Scikit-learn over the original Scikit-learn (training)




Speedups of Intel® Extension for Scikit-learn over the original Scikit-learn (inference)





# High-performance ML using accelerators (cont'd)

- Option 3: JAX 
  - A numpy library re-compiled for high-performance numerical computing
    - using XLA & JIT, on GPUs and TPUs
  - Key features
    - Automatic differentiation:
      - Forward and reverse mode AD of arbitrary numerical functions
      - grad, hessian, jacfwd, jacrev
    - Vectorization:
      - SIMD programming via automatic vectorization
      - vmap, pmap
    - JIT-compilation:
      - XLA is used to just-in-time(JIT)-compile and execute JAX programs
      - faster CPU code and transparent GPU and Cloud TPU acceleration
  - [sklearn-jax-kernel](#) is in an early stage.



**bit.ly/hpsk1\_02**

# Outline

- Learning Scikit-learn basics
  - High-level overview of scikit-learn libraries
  - Practical usage of scikit-learn
    - Deep learning and scikit-learn
    - Scikit-learn extension libraries
- High-performance machine learning using scikit-learn-*ish tools*
  - Overview of performance issues in machine learning
  - General performance tips and tricks in scikit-learn
  - Making the computation faster
  - Processing large datasets

# Training on Large Datasets

(where data is available upfront for processing)

- Challenges for large datasets
  - Most scikit-learn estimators works for in-memory arrays
  - Needs different dataframes
  - Needs to implement different algorithms
- Dask for Machine Learning
  - Using dask array, dask dataframe
- Machine Learning with vaex.ml
  - Wrappers to scikit-learn
  - Predictive models not implemented yet.
  - Implemented standard data transformers and KMeans algorithms
  - Using Vaex dataframe

**bit.ly/hpsk1\_03**

# Data in Out-of-core Learning



- Batch vs. Streaming
  - Batch: stored on disk or database, available upfront for processing
  - Streaming: all data streams at different time points, not all available upfront.
- Streaming plays a pivotal role in big data processing
  - Streaming architectures: Lambda vs. Kappa
  - Streaming technologies: Kafka, Flink, Storm, Cloud-based solutions on AWS, GCP, Azure
- Training-Serving Skew
  - Performance difference between training+testing and production
  - Possible causes:
    - Sourcing data from different pipelines for training and prediction
    - Processed in an ad-hoc manner with many short cuts
    - A feedback loop between your model and your algorithm.
  - Solution: to process both training and prediction data as part of the same pipeline

# Out-of-core learning using scikit-learn

- Feature extraction for a subset of data
  - Stateful feature extractor to build a “hash table”:
    - Must know the complete feature set known in advance
    - In-memory mapping from the string tokens to the feature indices
  - Stateless feature extractor using hashing tricks:
    - Preprocessing vectorizers having no ‘fit’
    - It creates reduced dimensionality hashes of data
- Incremental learning algorithm
  - Learning without seeing all instances at once
  - Scikit-learn supports through **partial\_fit** API
    - Classifiers: Naive Bayes, Perceptron, SGD, Passive-aggressive classifiers
    - Regressors: SGD, Perceptron, Passive-aggressive regressors
    - Clustering: Mini-batch K-means, Birch
    - Feature extraction: Dictionary learning, PCA, Latent Dirichlet Allocation
    - Preprocessing: Standard, MinMax, MaxAbs scalars
  - Mini-batch sizes may or may not influence results.

# Out-Of-Core Scikit-Learn Demo

- Classification of text documents

- Using Reuters-21578 benchmark dataset:
  - multi-class (90), multi-label
  - 7769 training documents and 3019 testing documents
  - provided by the UCI ML repository
- Binary classification between the "acq" class and all the others
  - "acq" was chosen as it is more or less evenly distributed in the Reuters # files

class name	nr of documents		mean number of
	train	test	words in train set
1: earn	: 2877	1087	104.4
2: acq	: 1650	719	150.1
3: money-fx	: 538	179	219.0
4: grain	: 1133	149	223.6

- Scikit-learn

- Data streaming with batches
- Create the vectorizer
- Classifiers with `partial_fit`

```
data_stream = stream_reuters_documents()
get_minibatch(data_stream, n_test_documents)

vectorizer = HashingVectorizer(decode_error='ignore',
                               n_features=2 ** 18,
                               alternate_sign=False)
X_train = vectorizer.transform(X_train_text)

cls.partial_fit(X_train, y_train, classes=all_classes)
```

**`bit.ly/hpsk1_04`**



# Key Takeaways

- Learning Scikit-learn means:
  - Understanding the machine learning workflows
  - Learning its API conventions
  - Familiarizing its documentations
  - Inspecting the example codes
  - Considering the scikit-learn extensive libraries
- Making scikit-learn more productive means:
  - Making sure everything work fine without performance optimization
  - Carefully profiling the code to address the performance bottleneck
  - Using Joblib with `n_jobs`
  - Trying RAPIDS, sklearnex, vaex.ml if needed.